

Programming of Super Computers

Assignment 2:

Parallel Debugging with TotalView

Isaías A. Comprés Ureña
compresu@in.tum.de

Prof. Michael Gerndt
gerndt@in.tum.de

20-11-2015

1 Introduction

Ensuring correctness is an important part of the software development process. In addition to correctness, the time required to produce results can also determine the usefulness of the software. Scientific software is often parallelized to produce results in less time. In the specific case of scientific software, the timeliness of the solution may be of equal importance as its correctness.

Unfortunately, while the parallelization process improves performance and allows an application to produce timely results, it also negatively affects the ability of developers to ensure correctness. In this second assignment, we will focus on a tool that helps us ensure the correctness of scientific software: the TotalView parallel debugger.

TotalView is offered by Rogue Wave Software. The Leibniz Compute Centre (LRZ) provides access to it (via a module) to all users of the SuperMUC. TotalView provides both CLI and GUI interfaces. The GUI allows the inspection of source code directly, control of threads and processes, memory and state inspection, etc. It supports C, C++ and Fortran applications as well as CUDA for NVIDIA GPU kernels in recent versions.

In this assignment, students will learn how to prepare and build a parallel application so that it can be analyzed with TotalView. The students will first have to understand common issues that arise when developing OpenMP and MPI applications. Afterwards, they will learn how to find, identify and fix errors in the application with the aid of the parallel debugger.

2 Submission Instructions and General Information

Your assignment submission for Programming of Supercomputers will consist of 1 part:

- A 5 to 10 minute video with the required comments described in each task.

There will be no extra files to be delivered in this assignment (only the video). Please refer to Section 2 of Assignment 1 for general instructions and recommended tools.

3 Understanding Parallel Programming Errors

Parallel programming is necessary to take advantage of supercomputing hardware. Due to the wide availability of multi-core CPUs, in recent years parallel programming is now also necessary to take advantage of most desktop and laptop PC hardware.

Going from single-threaded to parallel programming is difficult. It may require a complete redesign of the software: replacement of algorithms and data-structures, addition of locks, careful management of memory

and file descriptors, creation and tracking of threads and processes, etc. During these redesign activities, many common errors are introduced into the software. It is important to understand these common errors so that they can be avoided, and in the worst case identified, located and fixed.

In this task, do a short survey on the following common parallel programming errors (that affect correctness):

- Race Condition
- Deadlock
- Heisenbug (observer's effect)
- Floating-point arithmetic challenges
 - Comparisons
 - Definition of a zero and signed zeros
 - Cancellation or loss of significance
 - Amplification and error propagation

3.1 Required Video Commentary

Explain each of the above common errors. Mention which types of errors above are not exclusive to parallel programming. You are allowed to capture online resources in the video, such as illustrations, to help present your explanations.

4 Introduction to TotalView

TotalView is a parallel debugger offered by Rogue Wave Software. It is supported in several large compute centers in the world. With TotalView, developers can inspect source code, control threads and processes, inspect memory and state, etc. It supports C, C++ and Fortran applications, as well as multiple threads with OpenMP and multiple processes with MPI.

Before continuing with the subsequent tasks, take some time to learn the basics of TotalView. There are several online resources directly from Rogue Wave Software, as well as from some compute centers and universities. Here are some of them:

- Official TotalView Documentation from Rogue Wave Software
- TotalView at Leibniz Supercomputing Centre
- TotalView tutorial at Laurence Livermore National Labs

The Leibniz Compute Centre (LRZ) provides access to it (via a module) to all users of the SuperMUC. Some of the instructions in the LRZ's website (in the above list) are outdated, but the instructions on how to load and launch TotalView are correct.

The TotalView installation in SuperMUC seems to only work correctly in the "Phase 1" nodes, at the moment. Please make sure that you login to the "Phase 1" nodes and that you enable X11 forwarding and compression for your session:

```
ssh -YC <lrz-user>@sb.supermuc.lrz.de
```

You will need to use the IBM MPI and TotalView modules for this assignment. Make sure that you load them in your session:

```
module load totalview
module load mpi.ibm
```

The IBM MPI module is loaded by default, so it may not have to be loaded again. Make sure that you have the correct modules loaded by issuing the following command:

```
module list
```

and make sure that its output matches the following:

Currently Loaded Modulefiles:

```
1) admin/1.0  4) intel/15.0  7) mpi.ibm/1.4
2) tempdir/1.0 5) mkl/11.2   8) lrz/default
3) lrztools/1.0 6) poe/1.4   9) totalview/8.14
```

Once you have made sure that you logged in to the "Phase 1" login nodes of the SuperMUC, and that you have the correct modules to work with TotalView, you can now proceed to build your benchmark.

4.1 Preparing the Benchmark for TotalView

Refer to the instructions in assignment 1 and build 2 of the types of binaries: MPI and OpenMP. This can be done by building with the correct makefile options and then renaming the created binaries to avoid confusion. Additionally, make sure that the compiler flags are set to:

```
-g -O2
```

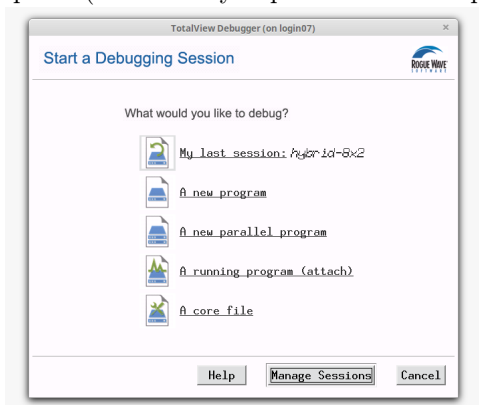
This is specified in the LRZ's TotalView instructions page. Please note that higher levels of optimizations will affect the ability of TotalView to match your source code, so please enable only the flags specified above. Rebuild your benchmark so that the new flags are now in effect.

4.1.1 Single Process Session Creation

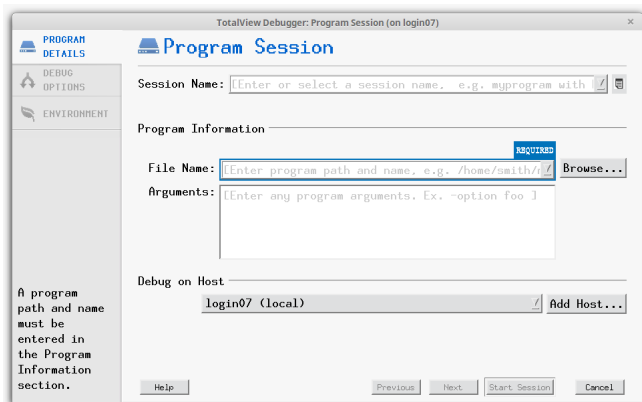
Once you have built the 2 binaries (MPI and OpenMP) and given them unique names, make sure that you are in the directory of the benchmark, and launch TotalView:

```
cd <benchmark-directory>
totalview &
```

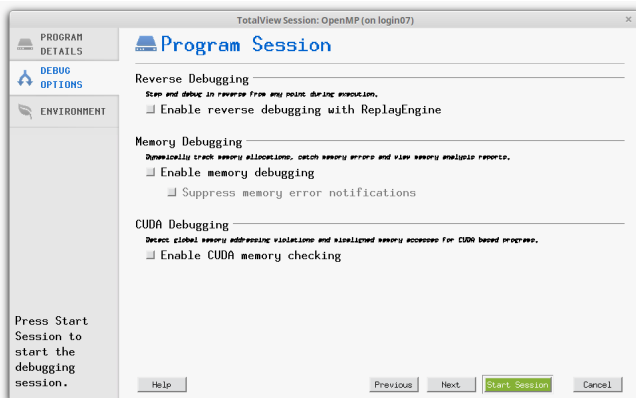
This will start the TotalView GUI. You should get the new windows created forwarded to your personal computer (make sure you passed the `-YC` options to ssh, as instructed before):



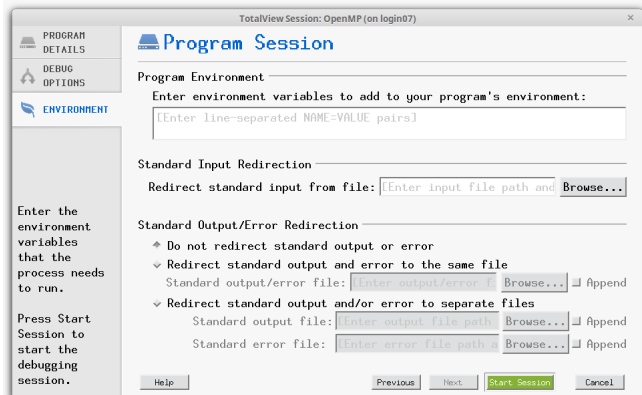
In the future, you can click the Manage Sessions button instead, to use, edit, duplicate or delete pre-existing sessions. Select the option A new program to create your first session:



Give the session name "OpenMP". Add the OpenMP version of the binary by clicking the Browse... button and using the file dialog. Note that the file is in the SuperMUC, and not in your local file system. You can leave the Arguments field empty, if your benchmark does not require any extra parameters to be launched. Click the Next button to continue:



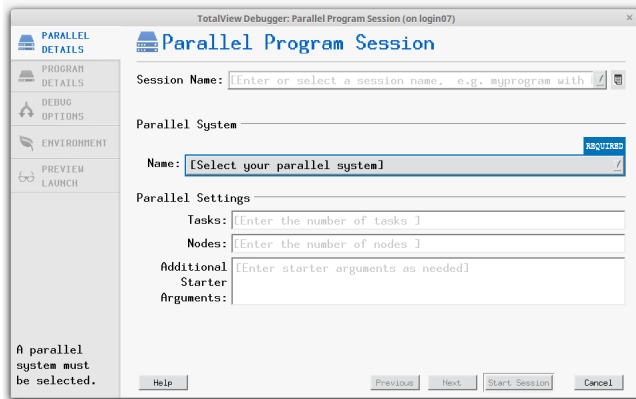
In the DEBUG OPTIONS screen, make sure that you do not enable Reverse Debugging since it currently crashes in the SuperMUC installation. You can enable Memory Debugging later if needed, at the cost of performance. Make sure that you do not enable CUDA Debugging, since the SuperMUC nodes do not have NVIDIA GPUs. Click the Next button to continue:



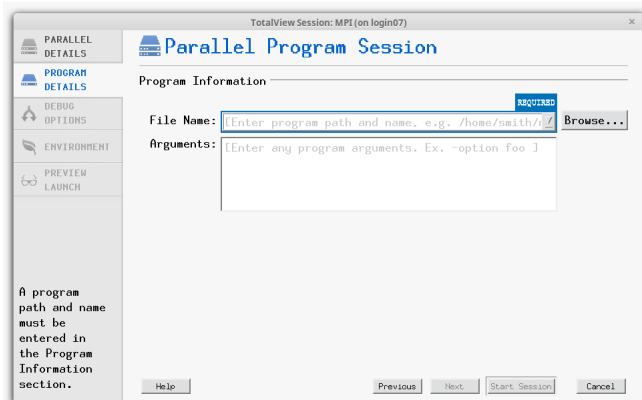
In the final dialog, you can skip all fields and simply click on Start Session to launch the program. The session will be saved and you can access it later via the session manager.

4.1.2 Parallel Session Creation

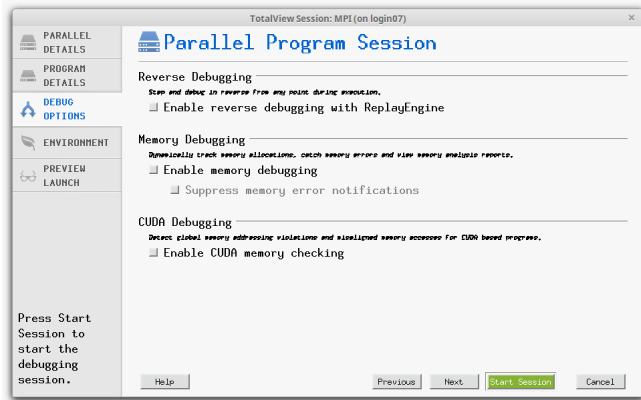
To set up a session for the MPI binary, you can go back to the session manager by using the menu File->Manage Sessions... from the main TotalView window, or using the welcome screen once you restart TotalView. This time, select A new parallel program:



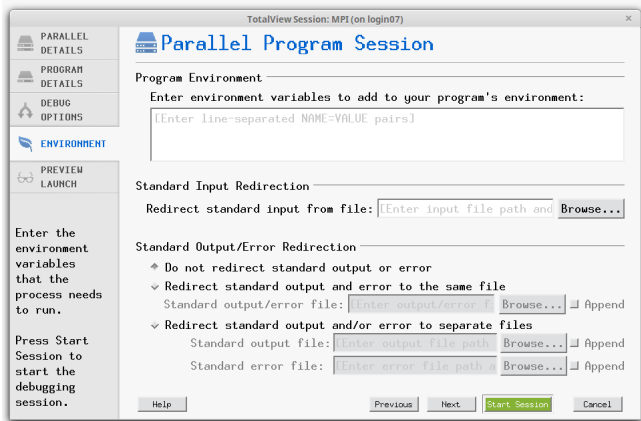
In this case, specify MPI as the name for the new session. In the SuperMUC, the Parallel System needs to be selected in the drop-down menu. Pick poe - Linux. Under Parallel Settings, set the Tasks (-procs) option to the required number of processes to run the benchmark in 2 nodes. Each of the thin nodes has 16 cores, so make sure that you check in your benchmark's documentation (refer to Assignment 1) that you pick the maximum number that can be run in 2 nodes. Set the Nodes (-nodes) field to 2.



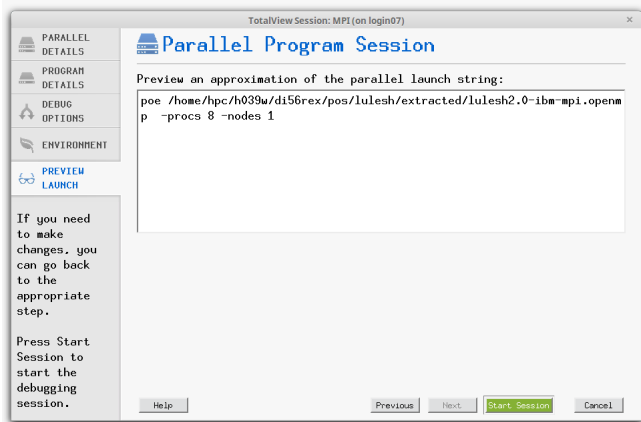
In the PROGRAM DETAILS tab, make sure to pick the correct binary for the MPI session by clicking in the Browse... button and using the file dialog. You can leave the Arguments field empty, in case your application requires no extra parameters.



The DEBUG OPTIONS screen is the same as for the single process case. Again, make sure that you do not enable Reverse Debugging since it currently crashes in the SuperMUC installation. You can enable Memory Debugging later if needed, at the cost of performance. Make sure that you do not enable CUDA Debugging, since the SuperMUC nodes do not have NVIDIA GPUs. Click the Next button to continue:

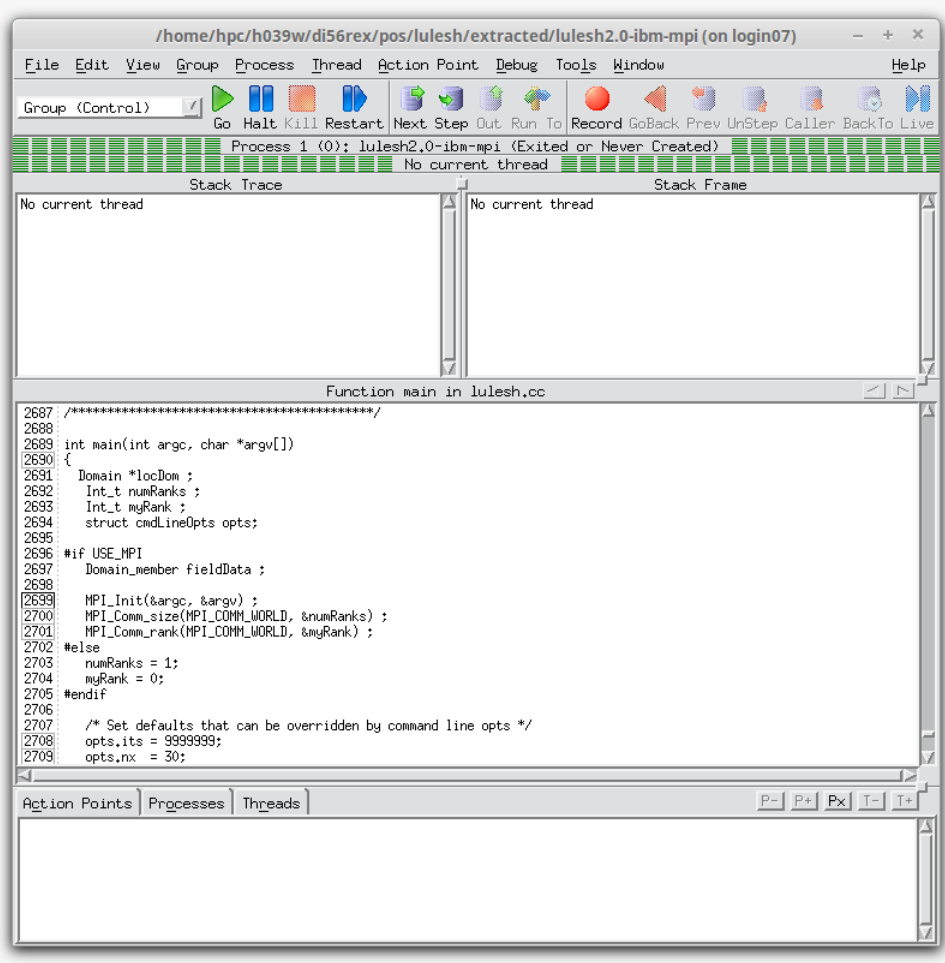


The ENVIRONMENT tab can be skipped. Click the Next button to continue:



The last tab is just a preview of the launch command. You can now click the Start Session button to launch the parallel application.

4.2 TotalView GUI



In this assignment, we will be working with the GUI version of TotalView. Make sure that you understand the following aspects of TotalView's GUI:

- Session Manager
- Root Window
- Process Window
 - Stack Trace Pane
 - Stack Frame Pane
 - Source Pane
 - Action Points, Processes, Threads Pane
- Variable Window

4.2.1 Required Video Commentary

Explain the purpose of each of the above listed TotalView GUI components. Show the relevant GUI components in the video as you explain it.

4.3 TotalView Basic Operations

Make sure that you learn how to do the following tasks from the tutorials and documentation of TotalView:

- Control execution
- Setting breakpoints
- Diving into functions
- View memory (variables and arrays)

4.3.1 Required Video Commentary

Perform the above operations with TotalView and capture them in the video. Give a short explanation on why these operations are important in a debugger.

5 OpenMP with TotalView

OpenMP uses a fork-join model of execution. The definitions of parallel regions, parallel loops, etc., are defined via OpenMP pragmas. Make sure to investigate the location of these pragmas in the benchmark.

Set any necessary breakpoints and then launch the single process session created in section 4.1.1. Use TotalView's GUI to iterate through the program and dive into a routine where a parallel region is entered and the new threads are forked by the OpenMP master thread.

5.1 Required Video Commentary

In the video, capture the before and after effect of the fork-join model, by navigating to a parallel region and looking at the `Threads` Pane in TotalView. Explain what the fork-join model is, along the way.

6 MPI with TotalView

MPI applications run with multiple processes and communicate using messages. Each process needs to call `MPI_Init` in order to set up its local data-structures and then be able to send messages and synchronize with other processes.

Launch the MPI session created in section 4.1.2. Use TotalView's `Source` Pane to set breakpoints at `MPI_Comm_size` and `MPI_Comm_rank`. Use the `Stack Frame` Pane to observe if each process stores the group size and its own rank.

Take some time to investigate the dominant routine of the benchmark. Identify one array or vector that is important for the computation. Use TotalView's Visualizer to visualize one or more arrays. Refer to the following page in the tutorial: Visualizing Array Data.

6.1 Required Video Commentary

Capture how the variables for rank and size change as the MPI application initializes and then sets its control variables. Mention the name of the relevant variables (for rank and size) and capture them at least 3 ranks in the video. State the name of the array you decided to visualize and capture its visualization in the video. It is not necessary to understand the actual meaning of the values in the array.