

Technische Universität München

Assignment 4: MPI Collectives and MPI-IO

Programming of Super Computers

Friedrich Menhorn, Benjamin R  th, Erik Wannerberg
Team 12

February 2, 2016



Contents

1. MPI Collectives

- 1.1 Optimizations
- 1.2 Performance Measurements

2. MPI Parallel IO

- 2.1 Data Sieving and 2-Phase IO
- 2.2 Scalability of original application
- 2.3 Impact of parallel IO

1. MPI Collectives

- 1.1 Optimizations
- 1.2 Performance Measurements

2. MPI Parallel IO

- 2.1 Data Sieving and 2-Phase IO
- 2.2 Scalability of original application
- 2.3 Impact of parallel IO

Optimizations

Sending the dimensions of the matrix to all processes:

Original code:

```
// send dimensions to all peers
if(rank == 0) {
    int i;
    for(i = 1; i < size; i++){
        MPI_Send(matrices_a_b_dimensions
                 , 4, MPI_INT, i, 0,
                 cartesian_grid_communicator
                 );
    }
} else {
    MPI_Recv(matrices_a_b_dimensions, 4,
             MPI_INT, 0, 0,
             cartesian_grid_communicator, &
             status);
}
```

With collective operation:

```
// send dimensions to all peers
MPI_Bcast(matrices_a_b_dimensions, 4,
          MPI_INT, 0,
          cartesian_grid_communicator);
```

Optimizations (contd.)

Distribute matrix blocks among all processes:

Original code:

```
// send a block to each process
if(rank == 0) {
    int i;
    for(i = 1; i < size; i++){
        MPI_Send(A_array, ...);
        MPI_Send(B_array, ...);
    }
    for(i = 0; i < A_local_block_size; i
        ++){
        A_local_block[i] = A_array[i];
    }
    for(i = 0; i < B_local_block_size; i
        ++){
        B_local_block[i] = B_array[i];
    }
} else {
    MPI_Recv(A_local_block, ...);
    MPI_Recv(B_local_block, ...);
}
```

With collective operation:

```
// send a block to each process
MPI_Scatter(A_array, A_local_block_size,
    MPI_DOUBLE, A_local_block,
    A_local_block_size, MPI_DOUBLE, 0,
    cartesian_grid_communicator);
MPI_Scatter(B_array, B_local_block_size,
    MPI_DOUBLE, B_local_block,
    B_local_block_size, MPI_DOUBLE, 0,
    cartesian_grid_communicator);
```

Optimizations (contd.)

Distribute matrix blocks among all processes and do initial rotation of blocks:

Original code:

```
// send a block to each process
...
// fix initial arrangements before the
// core algorithm starts
if(coordinates[0] != 0){
    MPI_Sendrecv_replace(A_local_block,
        ...);
}
if(coordinates[1] != 0){
    MPI_Sendrecv_replace(B_local_block,
        ...);
}
```

With collective operation:

```
// send a block to each process and fix
// initial arrangements before the
// core algorithm starts
int *send_count= (int *) malloc(...);
int *displs_A = (int *) malloc(...);
int *displs_B = (int *) malloc(...);
int displ_coord_A_col = ...
int displ_coord_B_row = ...
int displs_local_A = ...
int displs_local_B = ...

MPI_Gather(&displs_local_A, ...);
MPI_Gather(&displs_local_B, ...);

for(i = 0; i < size; i++){
    send_count[i] = A_local_block_size;
}

MPI_Scatterv(A_array, ...);
MPI_Scatterv(B_array, ...);
```

Optimizations (contd.)

Collect results after computation loop:

Original code:

```
// get C parts from other processes at
rank 0
if(rank == 0) {
    for(i = 0; i < A_local_block_rows *
        B_local_block_columns; i++){
        C_array[i] = C_local_block[i];
    }
    int i;
    for(i = 1; i < size; i++){
        MPI_Recv(...);
    }
} else {
    MPI_Send(...);
}
```

With collective operation:

```
// get C parts from other processes at
rank 0
int C_local_block_size =
    A_local_block_size;
MPI_Gather(C_local_block,
    C_local_block_size, MPI_DOUBLE,
    C_array, C_local_block_size,
    MPI_DOUBLE, 0,
    cartesian_grid_communicator);
```

Optimizations (contd.)

Expectations

- **Performance:**

Higher performance due to highly optimized MPI collectives (might also consider network topology etc.) vs. naive implementation via `MPI_Send/MPI_Recv`.

- **Scaling:**

Efficient implementation of Broadcast, Gather, Scatter use treelike distribution of values, $\mathcal{O}(n) \rightarrow \mathcal{O}(\log(n))$. This means we also benefit with respect to scalability.

Optimizations (contd.)

Readability and Maintainability

- ⊕ Sending dimensions with `MPI_Bcast`
- ⊕ Distributing matrix blocks `MPI_Scatter`
- ⊖ Distributing matrix blocks and initial alignment with `MPI_Scatterv`
 - introduces complicated offsets
 - hard to see what actually happens (distribution & rotation)
- ⊕ Collecting results using `MPI_Gather`

Performance Measurements

Show some measurements for falsification/verification of our guess

1. MPI Collectives

- 1.1 Optimizations
- 1.2 Performance Measurements

2. MPI Parallel IO

- 2.1 Data Sieving and 2-Phase IO
- 2.2 Scalability of original application
- 2.3 Impact of parallel IO

Data Sieving and 2-Phase IO

What is "Data Sieving" and "2-Phase IO"? How do they help improve IO performance?

Scalability of original application

- Was the original implementation scalable in terms of IO performance?
- Was the original implementation scalable in terms of RAM storage?

Impact of parallel IO

How much of the communication in the application was replaced with MPI-IO operations?